










IngestionJob
id: UUID kind: {arxiv_import, pdf_upload} state: JobState source_ref: str error: str [0..1] attempts: int

Paper
id: UUID source: {arxiv, upload} arxiv_id: str [0..1] {unique} arxiv_version: int [0..1] title: str authors: str [*] categories: str [*]

Analysis
id: UUID kind: {summary, synthesis, extraction} paper_ids: UUID [1..5] status: str result: JSON evidence: JSON (snapshots) prompt_version: str

Document
id: UUID sha256: str {unique} storage_key: str page_count: int extractor_version: str chunker_version: str embedding_model: str

Query
id: UUID question: str top_k: int paper_ids: UUID [*]

«enumeration»
JobState
queued, fetching, extracting, chunking, embedding, ready, failed


«enumeration»
AnswerStatus
answered, insufficient_evidence, error


Chunk
id: UUID ordinal: int page_start, page_end: int char_start, char_end: int token_count: int text: str embedding: vector(384) text_tsv: tsvector


Answer
id: UUID status: AnswerStatus reason: str [0..1] claims: JSON generation_model: str [0..1] prompt_version: str config: JSON latency_ms: float

AnswerEvidence
label: str (P1..Pn) rank: int score: float (cosine) chunk_id: UUID [0..1] paper_title: str page_start, page_end: int text: str (snapshot)


AnswerCitation
claim_index: int position: int label: str valid: bool







API request (synchronous)



Background job (JobRunner thread)

